

StressGuard: A Localized Edge AI System for Real-Time Stress Detection via High-Fidelity HRV Analysis

Arpita Hage - ah4384 & Urva Dhoriyani - ud2142
Columbia University
New York, NY, USA

Abstract

Chronic stress is a major contributor to cardiovascular disease and professional burnout. While wearable technology has advanced, cloud-dependent architectures introduce latency and privacy vulnerabilities. We present StressGuard, a localized, privacy-first edge system for real-time stress detection. Built on a Raspberry Pi 5 with a MAX30102 PPG sensor, the system extracts 14 clinical-grade Heart Rate Variability (HRV) features. Using an XGBoost classifier optimized via SMOTE and trained on the WESAD dataset, we achieved a classification accuracy of 78% with sub-200ms inference latency.

1 Introduction

Stress acts as a silent contributor to long-term health deterioration and mental burnout. The goal of StressGuard is to provide real-time, localized awareness to facilitate timely intervention. By performing all computation on the edge using a Raspberry Pi 5, we eliminate the need for cloud-based processing. Our approach addresses the critical technical gap between low-cost hardware and clinical-grade analysis. This localized compute node ensures data sovereignty and eliminates internet dependency for continuous monitoring. The system provides a privacy-first localized compute environment to protect sensitive health data. Real-time feedback is critical for respiratory recovery and acute stress management. Localized systems significantly reduce security risks associated with data transmission. The hardware utilizes accessible components to democratize health monitoring technology. Continuous monitoring allows for the identification of physiological trends over extended periods.

2 Physiological Background

2.1 The Autonomic Nervous System (ANS)

The human stress response is regulated by the ANS, which maintains internal homeostasis. The ANS consists of the Sympathetic Nervous System (SNS) and the Parasympathetic Nervous System (PNS). Stress induces SNS dominance, leading to the "fight-or-flight" physiological response. Conversely, the PNS governs the "rest-and-digest" phase, associated with higher variability. StressGuard extracts these delicate indicators directly from raw Photoplethysmogram (PPG) waveforms. Autonomic balance is quantified through the relationship between SNS and PNS activity. Sympathetic dominance is marked by elevated heart rates and decreased inter-beat variability. The system monitors the parasympathetic tone to identify calm physiological states. Physiological states like amusement and baseline are also part of the autonomic spectrum. The system provides a non-invasive window into the user's internal biological stress levels. Localized compute nodes allow for faster detection of

these autonomic shifts. This allows for immediate user intervention through breathing exercises or mindfulness.

2.2 Heart Rate Variability (HRV) as a Metric

HRV measures the variation in time between each heartbeat, known as inter-beat intervals (IBI). It serves as a non-invasive proxy for autonomic nervous system activity. High variability typically indicates a healthy, resilient state with high vagal tone. Low variability is an established clinical marker for psychological and physical stress. The system identifies R-peaks from raw PPG signals to calculate these intervals with high precision. Mathematical features are derived from the inter-beat interval time series. These metrics provide a quantitative assessment of the heart's rhythm stability. Clinical reference ranges are utilized to differentiate between stress and calm. Fourteen specific HRV metrics are utilized for the machine learning model's input. HRV is more informative than simple heart rate for understanding emotional state. The system's HRV analysis matches physiological expectations during live testing. The high-resolution polling ensures that the subtle differences in beat-to-beat timing are captured accurately.

3 System Architecture




Figure 1: System Architecture Diagram: Illustrating the three-layer pipeline from sensing to ML inference.

3.1 Sensing Layer

The MAX30102 Pulse Oximeter serves as the primary sensor, capturing raw PPG data at 100Hz. It utilizes integrated IR and Red channels

to detect blood volume changes at the fingertip. An MLX90614 sensor is optionally integrated via I2C for skin temperature monitoring. System temperature readings are also derived from the MAX30102 internal die temperature. The sensing layer communicates with the main unit through a stable 100Hz I2C hardware polling rate. Raw PPG waveforms are polled continuously to ensure no data points are missed. The I2C protocol allows for a modular sensor setup for future expansion. A 100Hz sample rate is required to maintain the fidelity of the HRV time series. Hardware stability was observed with no buffer overflows during extended sensing tests. Multimodal integration is planned for future iterations to include extra physiological markers. The sensor's low power profile makes it ideal for long-term wearable use cases.




Figure 2: Sensor Close-up: Physical layout of the MAX30102 sensor module and optical components.

3.2 Processing Layer

The Raspberry Pi 5 (8GB) acts as the central hub for all signal processing. It executes continuous polling and signal filtering to remove artifacts from the PPG stream. The system implements a configurable 15 to 30-second measurement window for feature extraction. During this phase, 14 specific time and frequency-domain metrics are calculated locally. A serialized ML model then processes these features to deliver an immediate classification. Feature extraction includes metrics like RMSSD, SDNN, and LF/HF ratio. The Pi 5's high RAM and processor speed allow for TFLite-optimized inference. A sliding window approach is used to maintain real-time

responsiveness. Processing involves mathematical consensus from the targeted feature set. The pipeline is designed to be modular for easy debugging and scaling. Data Persistence is handled via secure local CSV logging to analyze long-term trends. The system uses a dedicated thread for data acquisition to avoid polling latency during inference.

3.3 Output Layer

A live terminal interface provides real-time physiological feedback and stress status. The interface displays current Pulse (BPM), board temperature, and RMSSD values. Color-coded alerts notify the user of significant physiological changes. Secure local CSV logging is implemented for long-term health trend analysis. A Tkinter-based GUI is planned for future enhancements to include wellness tips. The terminal output renders the final status classification as Normal or Stress. Live status updates ensure the user is aware of their current autonomic state. The system provides immediate visual feedback to the user on the edge device. Output rendering is optimized to minimize total system latency. Future updates will include a mobile dashboard for data synchronization. Each alert is timestamped and recorded to the local file system for retrospective review.

4 Mathematical Foundations

4.1 Time-Domain Analysis

RMSSD is the primary indicator of parasympathetic activity, calculated from interval differences. SDNN represents the standard deviation of NN intervals over the entire window. pNN50 measures the proportion of RR intervals differing by more than 50ms. Mean HR provides a baseline of heart rate activity during the measurement period. These mathematical features map directly to established clinical stress thresholds. The system calculates inter-beat intervals (RR intervals) from the detected R-peaks. RMSSD specifically evaluates the short-term variation in heart rate rhythms. SDNN quantifies the overall heart rate variability during a recording session. pNN50 provides insight into the parasympathetic nervous system's responsiveness. All time-domain features are calculated entirely on the edge hardware. Mathematical consensus ensures that outliers do not trigger false stress alerts. The resulting feature vector is normalized before being passed to the classifier.

4.2 Frequency-Domain Analysis

Spectral analysis utilizes the LF/HF Ratio to quantify autonomic balance. The ratio is derived by integrating the Power Spectral Density (PSD) over specific bands. Low Frequency (LF) components generally reflect both sympathetic and parasympathetic activity. High Frequency (HF) components are primarily associated with parasympathetic activity. Elevated ratios strongly correlate with high-stress states and sympathetic dominance. The PSD is integrated between 0.04 Hz and 0.15 Hz for the LF band. The HF band integration occurs between 0.15 Hz and 0.4 Hz. Frequency domain features provide a deeper look into the SNS/PNS balance. These metrics are essential for clinical-grade HRV analysis on the edge. The LF/HF ratio is a top predictor in the stress classification model. Calculations are performed in real-time to provide immediate spectral feedback. The use of Fast Fourier Transform (FFT) enables efficient calculation of these power bands on the Pi hardware.

5 Machine Learning Methodology

5.1 Dataset: WESAD

The model was trained offline using the WESAD dataset. This dataset incorporates multi-modal signals from baseline, stress, and amusement states. Signals from the dataset were processed to mirror the features extracted by our hardware. The inclusion of diverse emotional states ensures the model can differentiate stress from arousal. Class imbalance in the dataset was addressed using the SMOTE technique. WESAD is a public benchmark for wearable stress and affect detection. The training pipeline used 15-fold Leave-One-Subject-Out Cross-Validation (LOSO-CV). Model training incorporates data from multiple physiological signals. The dataset allows the model to learn the mathematical nuances of different stress levels. Ground truth labels from the dataset were used to validate classifier accuracy. The training phase focuses on creating a generalized model for unseen subjects. The dataset's wrist-PPG segment was prioritized to align with the hardware's sensor modality.

5.2 Classification Pipeline

An XGBoost classifier was selected for its performance on categorical physiological data. A Random Forest variant was also tested for low memory footprint and high accuracy. The final classification is derived from the consensus of multiple decision trees. Serialization via joblib allows for rapid edge deployment. Top predictors identified include CVNN, HR_mean, and MedianNN. TFLite optimization was utilized to reduce model size for the Raspberry Pi 5. Inference mechanisms prioritize low memory usage and high speed. The classification result differentiates between Normal/Calm and Stress Alert. The model derives classification from the mathematical consensus of 14 features. Offline training ensures that the edge device only handles the lightweight inference task. Rapid deployment is achieved through pre-trained weights and optimized code. Hyperparameter tuning was conducted to maximize the F1-score across different user profiles.

6 Results and Performance

6.1 Model Efficacy

Leave-One-Subject-Out Cross-Validation yielded a binary accuracy of 76-78%. The best performing classifier was Logistic Regression at 78.07%. The XGBoost + SMOTE configuration achieved a stable 76.56%. These results are consistent with existing wrist-PPG HRV literature. Live system validation successfully differentiated CALM from STRESS ALERT states. The model accuracy reaches clinical-grade standards for daily management. Pulse and HRV readings are consistent with physiological expectations during testing. Hardware polling rates are stable, preventing data loss in high-stress detections. A binary classification provides a clear alert for the user to take action. The system shows robust performance across different subjects in the validation phase. Overall efficacy demonstrates the feasibility of edge-based clinical monitoring. The stability of the results across 15 subjects suggests high model reliability.

6.2 Hardware Benchmarks

The Raspberry Pi 5 maintained a stable 100Hz I2C hardware polling rate. No buffer overflows were observed during extended testing of the data acquisition pipeline. End-to-end inference latency, including feature extraction, was measured at sub-200ms. The measurement window completes in 15-30 seconds followed by immediate rendering. Hardware stability ensures robust performance for daily stress management. Latency benchmarks confirm a smooth real-time user experience. PPG polling begins at $T = 0s$ and the measurement window ends at $T = 15s$. Final terminal output is rendered within milliseconds of the window closing. The system consumes low power while maintaining high computational throughput. I2C pins provide a reliable interface for high-speed sensor data. Full latency benchmarking on the Pi 5 is currently in its final stages. Thermal monitoring shows that the Pi 5 operates within safe temperatures even during continuous inference.




Figure 3: Full System Assembly: The final hardware configuration including sensor mounts and the Raspberry Pi 5.

7 Conclusion

StressGuard balances clinical consistency with robust hardware stability. We successfully demonstrated 100% data sovereignty by performing all compute locally. The system provides reliable feedback for daily stress management with high accuracy. Our optimized pipeline provides a low-latency framework suitable for future clinical expansion. Future scope includes the integration of Electrodermal Activity (EDA) and mobile synchronization. The

system bridges the gap between affordable hardware and high-end health tech. Edge AI deployment proves to be a viable path for privacy-focused medical devices. Results are consistent with published literature in the field of PPG-based HRV. The project successfully delivered a working on-device stress detection prototype.

StressGuard serves as a modular foundation for multimodal physiological monitoring. The final outcomes provide a stable, low-latency framework for real-world stress tracking.